

### Part 3A

For NULL-pointer dereference, I basically mimicked what is done when creating the user stack by adding an inaccessible page starting at virtual address 0. To do this, in `exec` before allocating pages for the user code, I allocate an empty page that is mapped to virtual address 0x0 (using `allocuvm`) and then clear its user bit (`PTE_U`) to make it inaccessible to users. So if the user program tries to access any memory in this page, when the hardware does the translation, it will see that the PTE's user bit is cleared and will end up throwing a pagefault error.

To implement the memory protection functionality, I added the system calls `mprotect` and `munprotect`. In these functions, I leveraged the code in `xv6` to first get the arguments passed into the syscall (using `argint` and `argptr`). `Argptr` handles most of the error handling to check if the address is valid. The only thing I had to check myself was if it was page aligned. Once I got the arguments and error checked them, I used `myproc` to get the `pgdir` of the current process and then `walkpgdir` to get the page table entry of the specified address range. For each page table entry, I either clear the `PTE_W` bit (`mprotect`) or set it (`munprotect`). To make sure that the hardware knows of this change, I write to the `cr3`.

### Part3B

To create a kernel thread, I copied `fork()` and instead of allocating a new address space, I copied the pointer of the parent process to the new thread. Additionally, to ensure that when the thread runs it runs `fcn` and has access to the arguments, I set the trap frame instruction pointer (`eip`) to `fcn` and stack pointer (`esp`) to the top of the stack that was passed in and also copied the arguments and return address to the top. So `esp` ended up pointing to the return address. `Join()` is the same as `wait()` except for now I check that the parent and child process have the same `pgdir` to determine if it's a thread. I also no longer free up the address space. To keep track of the thread stack, I added an element to the `proc` struct and store the stack pointer there. I also added a thread count to the `proc` struct. It is initialized to 0 in `allocproc`, incremented by 1 in `clone`, and decremented by 1 in `join`. This is then used in `wait` to make sure that there are no more active threads for a process so that the address space can be freed.

For locks, I leveraged the code from wiki to add a fetch-and-add atomic operation in `x86.h`, which I use to implement the ticket spin lock from `OSTEP`.

For the thread user API, in `thread_create` I call `malloc()` to allocate a page of memory for the thread stack and then pass that in along with the user arguments to `clone()`. In `thread_join`, I call `join()` and if it returns a valid PID, I free up the stack pointer that it places in the argument that is passed in to it.

When a process' address space is grown, I added a lock acquire in `growproc()` to make sure the address space of a process and its children threads is synchronized. Other than a lock, I also added a loop through the process table to find all of the process' child threads and update their address space size (`p->sz`).