

For `kmalloc`, I basically reused all of the `umalloc.c` source code. Since I couldn't use `sbrk()` to request more memory for the kernel, I replaced that call with `kalloc()`. This means that when the kernel asks for more memory, `kmalloc` either gets already allocated memory from a free list or requests a whole page of memory from `kalloc()`. In order not to waste all of this page, the code breaks it down into `sizeof(Header)` chunks. So when the kernel calls `kmalloc`, it gets  $(n/\text{sizeof(Header)} + 1)$  chunks allocated to it. This is kept in a list of in-use chunks. When `kfree` is called, the chunks associated with that memory address are placed in the free list.

To implement `mmap`, I added another element to the process struct called `maps` which is just a linkedlist that contains all of the mappings for a specific process. I use `kmalloc` to allocate space for the linkedlist. A single node of the linkedlist contains elements that help keep track of the metadata such as length of mapping, mapping address, flags, fd, etc. Since we are only doing anonymous mappings without caring about the flags (getting implemented in project 5), the only two arguments I need to take into account are address and length. So when a process calls `mmap`, the first thing that I do is ensure that the inputs are valid: ensure `len > 0`. If it is  $> 0$ , then I round the length up to the nearest page. I do this since all memory blocks returned by `mmap` need to be page aligned. I then check the address hint. If the address hint is  $\geq \text{KERNBASE}$  and  $> p\text{->sz}$ , then the address I will try to map it to is `PGROUNDUP(addr_hint)`. If the hint does not satisfy those restrictions, then I by default will try to map it to an address base I created specifically for `mmap` called `MMAPBASE` (`0x40000000`). Once I have a `target_address` mapping, I use `walkpgdir` to check if the mappings from `target_address` to `target_address + length` already exist (check if it exists in the process `pgdir` and if the `PTE_P` flag is set). If any of them already exist, then it increases the `target_address` by length and repeats the previous step. It does this once it either finds a free chunk of memory that has not been mapped or it reaches `KERNBASE`. If it reaches `KERNBASE` then it returns an error. If it finds a free chunk of memory that can be mapped, then it allocates memory and maps it to `target_address` using `allocuvm()`. If this is successful, then it creates a new linkedlist node to store all of the metadata for this mapping and adds it to the process' `mmap` linked list, `p->mmaps`. `Munmap` is a bit more straightforward. All I do is iterate through the process' `mmaps` linked list to check if the address that is passed is a valid mapping. If it finds that address in the linked list, then it calls `deallocuvm()` to free that memory and unmap it from the process' address space. Once that is done, it then removes the mapping's node from the `mmaps` linkedlist and frees that node using `kfree()`. One list thing I added is code in `allocproc()` to initialize the process' `mmaps` linked list to empty (null) and code in `freevm()` to destroy the process' `mmaps` linkedlist by calling `kfree()` for all the nodes in the list. There is no need for me to free and unmap the `mmap` mappings since that is already done in `freevm()`.